

The Enemy of My Enemy: How Competition Mitigates Social Dilemmas*

Alessandro Stringhi †

Andrea Albertazzi ‡

Prague University of Economics and Business

IMT School for Advanced Studies Lucca

Sara Gil-Gallen §

ISTC, Italian National Research Council

December 10, 2025

Abstract

We study whether competition between groups that yields no material rewards fosters within-group cooperation. In a laboratory experiment, pairs of subjects played an indefinitely repeated Prisoner’s Dilemma either in isolation or in a tournament against another pair, where winning conferred no monetary payoff. Competition increased cooperation, and the effect strengthened over time as subjects gained experience. Exploiting the binary action space of the repeated game, we estimate participants’ strategies and find that competition reallocates play away from strategies that start with defection toward strategies that start with cooperation. This reallocation results in a shift from Always Defect to the least risky cooperative strategy, Grim. We complement the experiment with a simple tournament-based model that shows how non-monetary competition can shift equilibrium selection toward cooperative strategies.

Key words: *Competition; Cooperation; Prisoner’s Dilemma; Repeated game.*

JEL code: C73, C92, D81.

*We thank Giuseppe Attanasi, Pablo Brañas-Garza, Nikolaos Georgantzis, Salvatore Nunnari, Simon Weidenholzer, and participants to the following conferences: IAREP-SABE 2023, EEA-ESEM 2022, ASFEE 2022, IX Doctoral workshop UAB 2021, and Rede3c seminars 2021. This research has benefited from the support of the French Agence Nationale de la Recherche (ANR) under grant ANR-18-CE26-0018 (project GRICRIS).

†Prague University of Economics and Business, *e-mail:* alessandro.stringhi@vse.cz.

‡IMT School for Advanced Studies Lucca, *e-mail:* andrea.albertazzi@imtlucca.it.

§Institute of Cognitive Sciences and Technologies, Italian National Research Council, *e-mail:* sara.gilgallen@cnr.it.

1 Introduction

Theoretical and empirical work shows that cooperation among group members is crucial to the performance of a wide range of institutions (e.g., Alchian & Demsetz, 1972; Fehr & Gächter, 2000; Hamilton, Nickerson, & Owan, 2003; Holmstrom, 1982; Ostrom, 1990). However, the conflict between selfish and cooperative choices endures in many social and economic interactions, presenting a trade-off between individual and collective interests.

Experimental literature has documented that intergroup competition represents a valuable intervention to mitigate this social dilemma and promote cooperation, but it is typically implemented in conjunction with financial incentives. This confounding makes it difficult to identify the effect of competition *per se*, as opposed to that of monetary rewards. At the same time, many real-world interventions in firms, schools, and the public sector introduce competition among groups without altering monetary incentives, for example through public rankings that compare departments, units, or branches. Therefore, understanding whether competition *per se* can foster cooperation remains of critical importance for both experimentalists and organizations seeking to enhance cooperation through cost-effective instruments.

In this paper, we investigate whether intergroup competition promotes within-group cooperation in the absence of material incentives. Our goal is not only to document changes in average cooperation but also to identify how competition reshapes strategic behavior. To this end, we run a laboratory experiment in which pairs of subjects play an indefinitely repeated Prisoner’s Dilemma game (PD). We compare a baseline to a treatment that introduces a tournament between pairs while holding stage-game payoffs fixed. Participants’ earnings depend only on outcomes in their own game, and winning the tournament yields no monetary reward. Thus, this design cleanly identifies the effect of competition absent material prizes.

We find that first-round (all-rounds) cooperation increases on average by $\approx 42\%$ ($\approx 15\%$) in the tournament compared to the control condition, and this effect unfolds and becomes significant as subjects gain experience. The result is driven by a shift in strategy adoption. In the baseline, the most common strategy followed by subjects prescribes unconditional defection, Always Defect. By contrast, in the treatment, a substantial share of subjects adopt Grim, the least risky cooperative strategy. Competition appears to foster a cooperative norm under which players mutually cooperate, and deviations are met with harsh punishment. These differences reveal that competition alone alters behavior along dimensions not visible in aggregate cooperation rates. By exploiting the binary and history-dependent structure of the repeated PD, our design, along with the strategy estimation, offers a key methodological contribution to the study of intergroup competition and cooperation.

Intergroup competition is pervasive in organizational, academic, and social contexts,

and experimental evidence consistently finds that it enhances cooperation within groups. However, most existing studies pair competition with monetary rewards (e.g., Abbink, Brandts, Herrmann, & Orzen, 2010; Y.-Y. Chen, 2020; Markusen, Reuben, & Tyran, 2014; Puurtinen & Mappes, 2009; Reuben & Tyran, 2010), making it difficult to disentangle competition from the direct effects of material incentives.¹ Moreover, the literature has primarily relied on public goods games (e.g., Augenblick & Cunha, 2015; Burton-Chellew, Ross-Gillespie, & West, 2010), where cooperation is proxied by average contributions, a coarse metric that offers limited insight into the underlying strategies that sustain cooperation. Relatively little is known about whether a competitive environment without monetary stakes fosters cooperation, or how such competition affects the behavioral patterns that lie behind similar aggregate cooperation rates. This paper aims to address both questions.

As a first contribution, we show that intergroup competition, absent material incentives, fosters within-group cooperation. This contrasts with much of the experimental literature, which employs monetary rewards, thereby confounding competition per se with monetary incentive effects. To the best of our knowledge, Cárdenas and Mantilla (2015) and Tan and Bolle (2007) are the only experimental studies that examine the effect of intergroup competition on cooperation. Using public goods games, both studies find that cooperation is higher when winning is materially incentivized, and Tan and Bolle (2007) also documents a positive effect of relative-performance feedback. While these studies share important elements with our setting, our design combines three features that have not been studied jointly: (i) it includes a condition without relative information, allowing us to identify the effect of competition per se; (ii) it employs an indefinitely repeated Prisoner's Dilemma rather than a public goods game; and (iii) it implements indefinite repetition, thereby avoiding endgame effects. Thus, we isolate competition from material incentives and, exploiting the PD's binary, history-dependent structure, uncover individual strategies.

Therefore, our second contribution is that we capture how competition changes strategic behavior by employing the indefinitely repeated PD. This methodological choice is particularly advantageous, as its binary decision framework (*Cooperate* or *Defect*) clearly represents participants' intentions. The decision to cooperate unequivocally provides a round-by-round indicator of cooperative behavior. By contrast, in public goods games, the broader choice set available to participants introduces potential ambiguity in interpreting individual contributions. For instance, a positive contribution below the group's average may be interpreted as free riding relative to peers, even though it is cooperative relative to a non-cooperative benchmark. Our use of the repeated PD thus offers a clean mapping from observed actions to cooperation and, together with strategy estimation, allows us

¹A related body of research explores the role of non-monetary incentives such as relative performance information (Schnieder, 2022). However, it focuses on individual effort, performance, and sabotage. In contrast, our paper focuses on cooperation.

to recover the dynamic rules sustaining cooperation rather than only aggregate levels. Estimating strategies from actual choices is not a trivial task for two main reasons. First, the number of possible strategies is virtually infinite (Fudenberg & Maskin, 1986). Second, while each choice is conditional on a specific history, we only observe one actual choice and not what subjects would have done at other decision stages. Our approach involves the use of a finite mixture model to estimate the proportion of participants in each treatment who employ a specific strategy. This widely used method provides insights into the decision-making processes that drive cooperation, offering a more nuanced understanding than studies that use different designs and focus on aggregate outcomes.

We complement the experiment with a theoretical analysis in which we develop a simple tournament-based model that rationalizes the observed reallocation of strategies. The model shows how introducing an intergroup tournament effectively lowers the discount-factor thresholds under which cooperation and Grim strategies can be sustained in equilibrium. We posit that players receive hedonic utility from winning the tournament, a parsimonious assumption about the utility function that applies to agents who are otherwise purely self-interested. Although other mechanisms, such as group identity or social framing (Bacharach, 1999; Ellingsen, Johannesson, Mollerstrom, & Munkhammar, 2012), can likewise foster cooperation, the objective of our framework is to demonstrate that the tournament itself can generate the observed shift toward cooperative trigger strategies, in line with the experimental evidence, and that this effect does not depend critically on any particular specification of the utility function. In doing so, we contribute to the theoretical literature on intergroup competition and repeated interaction, which has typically relied on either group-contingent social preferences, identity-based utilities, or material incentives (Cason, Lau, & Mui, 2019; Cheikbossian, 2012; B. Chen, Li, & Liu, 2025). Identity-based and social-preference approaches effectively modify the Prisoner's Dilemma payoff matrix by allowing others' material payoffs to enter utility, while contest models change material incentives by awarding additional monetary payoffs for winning. In contrast, our model leaves the stage-game monetary payoffs of the Prisoner's Dilemma unchanged and instead introduces a small non-monetary reward at the level of the intergroup contest. Formally, we treat the probability of winning the tournament as an increasing function of a pair's expected total payoffs for any fixed distribution of opponents' strategies. The resulting analysis is therefore partial-equilibrium but sufficient for our purpose, as it shows how, for a given competitive environment, the tournament shifts the relative sustainability and risk-dominance of Grim versus Always Defect, thereby rationalizing the observed reallocation of strategies.

Our findings show that intergroup competition, without monetary incentives, can be a cost-effective means of fostering cooperation in environments where individual and collective interests collide. In our setting, introducing a simple tournament between groups, while leaving material payoffs unchanged, increases cooperation and shifts play toward

more self-enforcing cooperative paths. This result is consistent with a mechanism in which a modest hedonic value of winning can substitute for high-powered financial incentives. Thus, we show that such tournaments can be particularly attractive for organizations that face budgetary or equity constraints, yet seek to enhance welfare in settings where cooperation is socially beneficial. At the same time, our design abstracts from potential negative externalities of competition, such as sabotage. Thus, our conclusions may not extend to contexts where groups can directly influence rival groups.

The remainder of this paper is organized as follows: Section 2 introduces the game and describes the experimental design; Section 3 reports the results; Section 4 presents the theoretical analysis of the game; and Section 5 concludes.

2 Methods

2.1 Experimental Design

This study implements a between-subjects design in which pairs of subjects play an indefinitely repeated PD game. The game is based on one of the treatments from Dal Bó and Fréchette (2011), with the individual payoff matrix of the stage game represented in Table 1. For the remainder of this paper, we will use the term round to refer to the stage where subjects make decisions. Players can choose between two actions at each round: *Cooperate* or *Defect*.² At the end of the round, there is a fixed and known probability $\delta = 0.75$ (continuation probability) that the game will continue, and the participant will play with the same partner in the next round. We refer to the series of consecutive rounds played with the same partner as a supergame.

	<i>Cooperate</i>	<i>Defect</i>
<i>Cooperate</i>	32, 32	12, 50
<i>Defect</i>	50, 12	25, 25

Table 1 Payoffs of the stage game represented in Experimental Currency Units (ECU).

At the end of each round, every player receives feedback about the action taken by their partner and the resulting outcome. A history box summarizing the actions and payoffs of both players in previous rounds is displayed on the screen and stored until the end of the supergame. When a supergame ends, new pairs are randomly formed, and a new supergame with the same rules begins. The participants have 50 minutes to play, and their earnings are computed as the cumulative sum of the individual payoffs of each round they played. We refer to the subjects who play using this set of rules as the *Control* group.

²To prevent unintended framing effects, the actions in the experiment were labeled as Action 1 and Action 2, respectively.

Tournament In the treatment group (henceforth *Tournament*), the rules of the game are identical to those of *Control*, with a single exception: the two players are competing with another pair of subjects. At the beginning of each supergame, two pairs are randomly matched. The pair that accumulates more points (the sum of both players' individual payoffs) by the end of the supergame is declared the winner. Because cooperation always yields more points to the pair, the more the two subjects cooperate, the greater their likelihood of winning the tournament. The outcome of the competition (win, loss, or tie) is displayed at the end of each supergame, but no information is given regarding the points achieved by the other pair. It is essential to note that winning the competition does not result in an additional monetary payoff, and participants are explicitly informed of this. If additional economic incentives were provided to the winners, it would not be possible to disentangle the effects of competition and monetary prizes, as the additional monetary rewards would virtually increase the stage-game payoff associated with cooperation.

2.2 Experimental Procedure

We recruited 94 participants (46 in the control group and 48 in the treatment group) from the subjects' pool of the University of Côte d'Azur (Nice, France) using ORSEE (Greiner, 2015). The subject pool included students from various disciplines. The experiment was programmed using zTree (Fischbacher, 2007) and conducted at the Laboratoire d'Économie Expérimentale de Nice (LEEN) in September 2020. The payoffs are expressed in Experimental Currency Units (ECU), and at the end of the experiment, participants were paid €0.50 for 100 ECU earned during the experiment. The average payment was €21.42, including a €5 show-up fee, and the experimental sessions lasted, on average, 75 minutes. We conducted a total of six sessions evenly distributed across treatments, and each participant played in one of the two treatments only. At the end of the experiment, participants completed a brief questionnaire in which they self-reported their socio-demographics, generalized trust, and risk aversion.³ Table C.1 in Appendix C shows that treatment randomization is balanced with respect to variables elicited in the final questionnaire.

3 Results

The primary objective of this study is to show that introducing a competitive environment that bears no additional economic rewards is sufficient to foster cooperation. The first part of this section reports evidence on the treatment effect, while in the second part, we provide results on the strategies estimation.⁴

³See the questionnaire in Appendix B.

⁴In one session of *Control*, we encountered a problem as subjects continued playing the game even after the 50-minute mark. For this reason, in that session, we only use observations that were played

3.1 Treatment effect

The left panel of Figure 1 shows cooperation rates in the first round of each supergame, a standard practice in the repeated-games literature. First-round choices are comparable across supergames with random lengths, avoid complications from history-dependent behavior in later rounds, and cleanly measure the willingness to start cooperating (see Dal Bó & Fréchette, 2011, 2018, for an in-depth discussion). Our data show a sizable treatment effect: average first-round cooperation rates are 38.74% and 54.92% in *Control* and *Tournament*, respectively. Introducing a competitive environment thus increased first-round cooperation by 16.18 percentage points ($p = 0.033$).⁵ The figure shows that the treatment effect emerges and increases in magnitude as subjects gain experience. Consistent with this, first-round cooperation in *Control* declines significantly over supergames ($p = 0.007$), whereas in the tournament we do not detect any systematic time trend ($p = 0.885$). This difference in time trends across experimental conditions is weakly significant ($p = 0.073$). In the right panel of Figure 1, we plot cooperation rates using all rounds within each supergame. Even though we do not detect systematic differences in supergame length across treatments (Mann-Whitney U test: $p = 0.1727$), these all-round comparisons should be interpreted with caution, as later-round actions are strongly history-dependent and yield multiple non-independent observations per supergame. As expected, treatment differences are attenuated once later rounds are included. Although the figure suggests higher cooperation in the *Tournament* condition toward the end of the experiment, the overall differences are not statistically significant when all rounds are considered. Overall cooperation rates are 31.36% in *Control* and 36.17% in *Tournament*, with no significant difference in levels ($p = 0.157$) or in time trends across treatments ($p = 0.353$).

Table 2 reports mean cooperation rates by treatment, separately for first rounds (left panel) and all rounds (right panel). Within each panel, columns Q1-Q4 isolate choices from the first, second, third, and fourth quartiles of the overall sequence of decisions, respectively. This aligns observations by their position in the overall sequence of play, rather than by supergame number, which can vary in length. The bottom rows display the treatment difference and the associated p-value for each comparison. The estimates indicate that the effect of competition on cooperation is consistently positive, negligible in the early stages of the experiment, but sizable and statistically significant in the final quartiles, where subjects have gained experience and behavior is likely to have stabilized. The treatment difference in the last quartile of decisions is more than 21 percentage points (+61%) in the first rounds and almost 10 (+34%) when all rounds are considered.

up to supergame 29 (the maximum supergame number played in other sessions). This ensures a similar number of total decisions across sessions and treatments, allowing for a fair comparison. Cooperation in those “extra” supergames further decreased over time, thus not undermining our results.

⁵Throughout the paper, if not otherwise stated, statistical significance is assessed by estimating average marginal effects from mixed-effects probit models with standard errors clustered at the participant level where cooperation is regressed against the relevant variables.

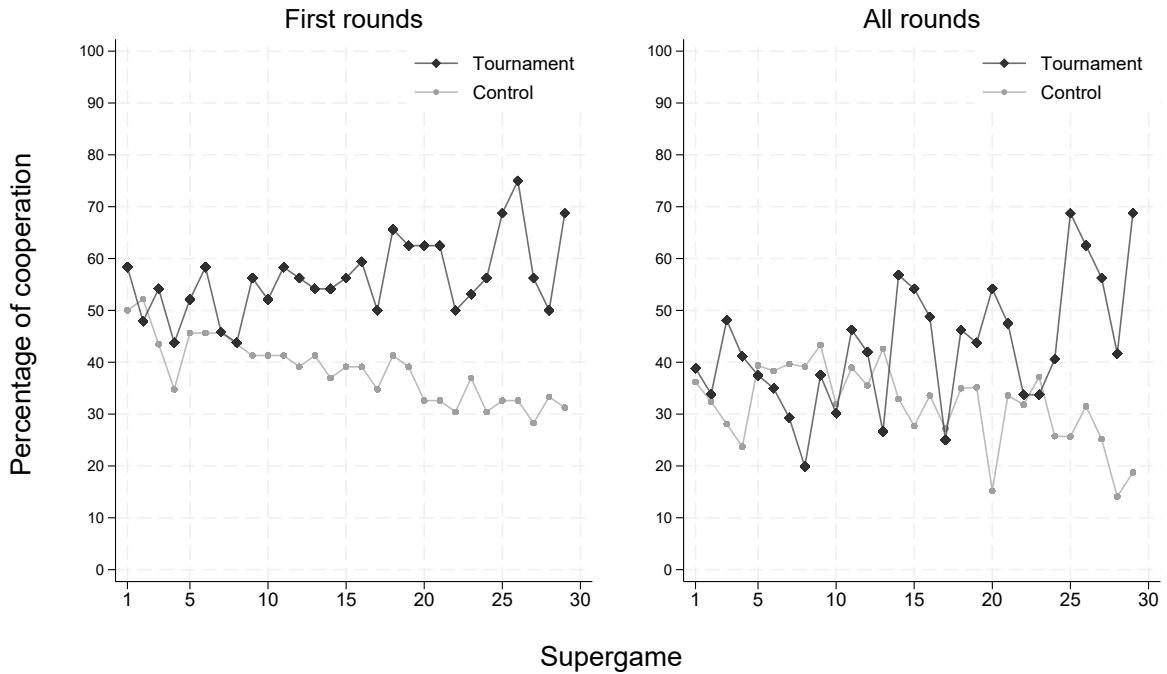


Figure 1 Percentage of cooperation by supergame. The left panel shows cooperation rates in the first rounds of each supergame. The right panel uses all observations.

	First rounds					All rounds				
	All	Q1	Q2	Q3	Q4	All	Q1	Q2	Q3	Q4
Control	38.74	47.02	40.10	36.56	35.71	31.36	32.79	32.43	31.14	29.26
Tournament	54.92	51.56	54.02	56.25	57.42	36.17	37.84	33.24	33.81	39.10
diff	16.18	4.54	13.92	19.69	21.71	4.81	5.05	0.81	2.67	9.84
p-value	0.033	0.513	0.302	0.075	0.003	0.157	0.403	0.489	0.436	0.027

Table 2 Percentage of cooperation in first rounds (left panel) and all rounds (right panel). Q1-Q4 shows cooperation rates in each quartile of the decision distribution.

The evidence presented so far highlights a sizable treatment effect that develops over time and consolidates in the last parts of the experiment. However, average treatment effects may conceal systematic changes in the composition of outcomes. In Figure 2 we display the distribution of stage-game outcomes by treatment, separately for first rounds (left column) and all rounds (right column), and for the full sample (top row) versus only the last quartile of decisions (bottom row). In every panel, the *Tournament* condition features a higher fraction of mutual cooperation (C,C) and a lower fraction of mutual defection (D,D) than the *Control* condition. The differences are especially pronounced in the last quartile of decisions (bottom row). Pair-level probit estimates indicate that the treatment significantly increases the probability of mutual cooperation in all panels except the top-right one, where the estimated effect is positive but not statistically significant at any conventional level (top-left panel: $p = 0.002$; top-right panel: $p = 0.403$; bottom-left panel: $p = 0.000$; bottom-right panel: $p = 0.006$).⁶

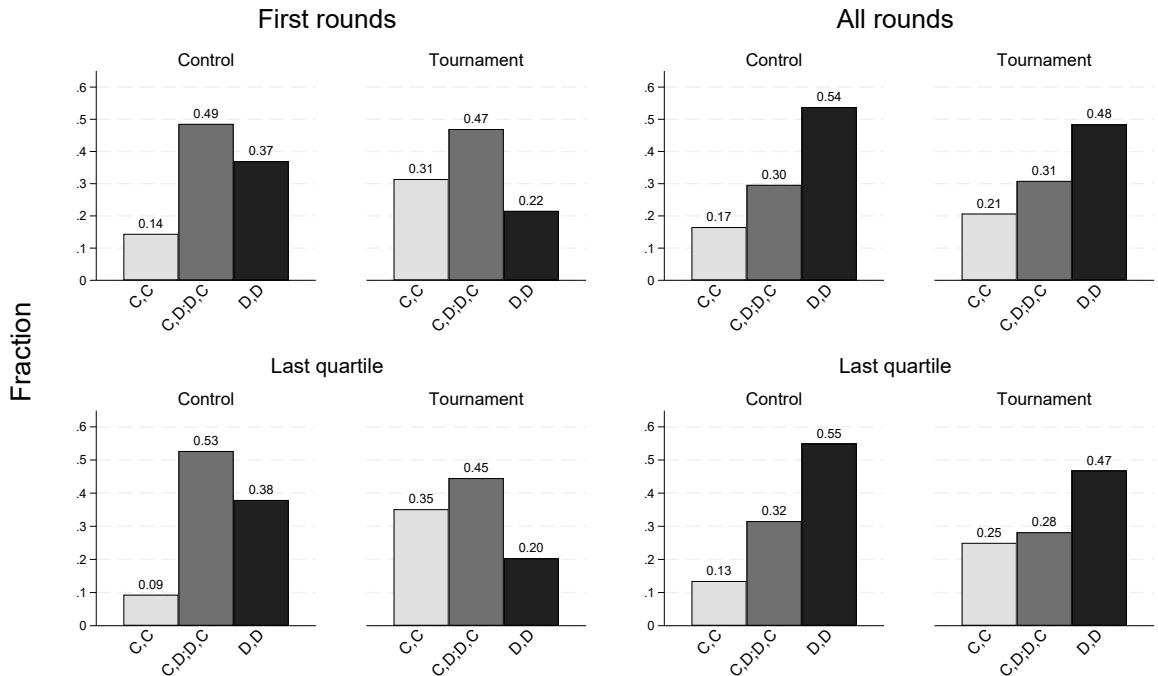


Figure 2 Outcomes by treatment. Pictures in the left column show the fraction of outcomes for only the first rounds of each supergame. The pictures in the second column use all rounds of every supergame. The first row uses choices from the entire experimental session, while the bottom row only uses decisions from the last quartile. The outcome in each round can be either mutual cooperation (C,C), one cooperates and the other defects (C,D or D,C), or mutual defection (D,D).

The pattern in Figure 2 suggests that the competition between groups primarily shifts mass from mutual defection toward mutual cooperation. Across panels, the fraction of

⁶Statistical significance is assessed using average marginal effects from probit specifications estimated at the pair level, with standard errors clustered at the session level.

D,D outcomes is systematically lower in *Tournament* than in *Control*, while the share of mixed outcomes C,D or D,C remains relatively stable. This suggests that the treatment does not merely induce one-sided deviations from defection, but rather fosters coordinated moves toward mutual cooperation.

In Figure 3, we visually investigate the role of previous outcome realizations on cooperation. The figure reports the probability of cooperation conditional on the previous realized outcome within a dyad, excluding first rounds so that comparisons reflect within-relationship dynamics rather than initial beliefs. We acknowledge that this represents a naive analysis of the mechanisms underlying the treatment effect, as it overlooks the within- and between-pair histories of choices. However, we believe it provides important visual evidence that motivates our strategy estimation presented in the next section.

Figure 3 shows a pattern more consistent with grim-like behavior in *Tournament*: higher cooperation rates after C,C , and lower cooperation rates after C,D , and D,D . These conditional histories visually show that competition increases cooperation after mutual cooperation, while leaving responses to deviations at least as punitive as those in *Control*. This pattern is consistent with a focal cooperative standard that prescribes playing C to remain on the path and is reinforced by strict deterrence, which punishes deviations from it. However, conditional frequencies cannot tell whether this operates through a reallocation across decision rules or changes in error rates within rules. This motivates the structural analysis we provide in Section 3.2, where we estimate a finite mixture model to quantify how competition reshapes the strategy mix.

3.2 Strategies

This section exploits the dependency of choices between and within supergames by estimating the strategies that participants used in the experiment via finite mixture models. This practice was pioneered by Dal Bó and Fréchette (2011) and later adopted and refined in subsequent experimental work (e.g., Aoyagi, Bhaskar, & Fréchette, 2019; Arechar, Dreber, Fudenberg, & Rand, 2017; Bigoni, Casari, Skrzypacz, & Spagnolo, 2015; Breitmoser, 2015; Camera, Casari, & Bigoni, 2012; Dal Bó & Fréchette, 2018; Dvorak, 2023; Fréchette & Yuksel, 2017; Fudenberg, Rand, & Dreber, 2012; Jones, 2014; Romero & Rosokha, 2023; Vespa, 2020).

We estimate strategies with a finite mixture model following Dvorak (2023).⁷ We start from a rich set of candidate strategies and assume that each individual is a latent type who follows one of these strategies, subject to random trembles.⁸ Given the observed action histories, we estimate the model by maximum likelihood and obtain, for each treatment,

⁷Please refer to the paper for details on the estimation procedure.

⁸Aligning with the literature, we assume a common probability of mistakes within each treatment rather than strategy-specific trembles. In Table D.2, we provide estimates using the selected set of strategies presented in the main analysis, relaxing this assumption. Our main conclusions do not vary.

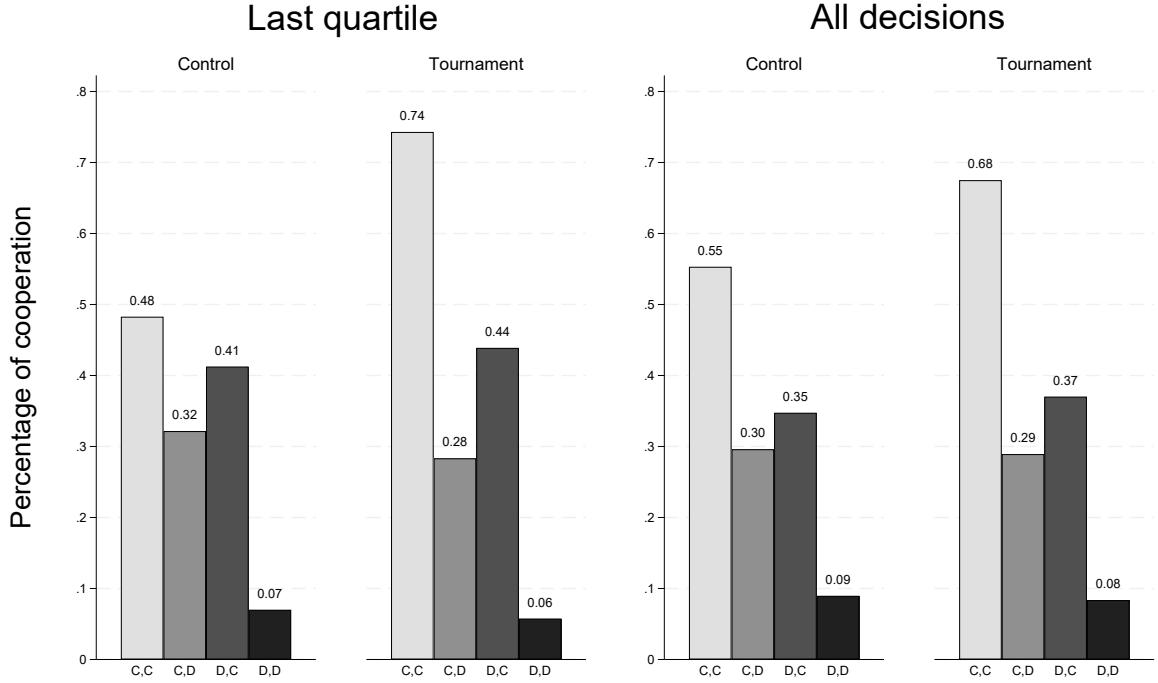


Figure 3 Percentage of cooperation conditional on the previous outcome. Only within supergame histories are considered. Hence, the first round of each supergame is excluded.

an estimated distribution over strategy types and an error rate. Because the space of possible strategies is virtually infinite, we restrict our attention to strategies that previous work has identified as relevant in indefinitely repeated games (Aoyagi, Fréchette, & Yuksel, 2024; Dal Bó & Fréchette, 2011, 2018; Fudenberg et al., 2012; Romero & Rosokha, 2023), which yields an initial set of 17 pure strategies.⁹ Following the literature, we focus this part of the analysis on later supergames (20–29), where behavior is likely to have stabilized, as suggested by the evidence provided above. When this was not possible, we attempted to maintain a balanced distribution of later decisions (interactions) across treatments. In practice, in both Control and Tournament, approximately the last quarter of decisions is used for estimation.¹⁰ Finally, because an overly rich candidate set can induce instability in the classification of observationally similar types (e.g., Grim vs Always Cooperate), we use the Bayesian Information Criterion (BIC) to select, from the initial set, a parsimonious subset of strategies to estimate.

The surviving strategies are: always defect (AD); Tit-for-Tat (TFT), that starts cooperating and then mimics the opponent’s choice in the previous round; Suspicious Tit-for-Tat (STFT), that starts defecting and then plays as TFT; Grim that cooperates until a defec-

⁹In Appendix D we describe the full set of candidate strategies, and in Table D.1 we provide estimates using this initial set. Our main conclusions remain unchanged.

¹⁰In Appendix D.2, we report results from a robustness check in which strategies are estimated using all available observations (supergames 1–29). The resulting estimates are qualitatively similar to those reported in the main analysis and do not alter our main conclusions.

tion from the opponent, then defects forever; Grim2 and Grim3 that are lenient version of Grim where defection is triggered by 2 or 3 consecutive defections of the opponent respectively; and T2 and T3 cooperate for 2 and 3 rounds respectively, and then defect forever. We then estimate a joint finite-mixture model allowing for treatment-specific mixture weights and trembles. Results are reported in Table 3.

	<i>Control</i>	<i>Tournament</i>
AD	0.380*** (0.078)	0.233*** (0.066)
TFT	0.159*** (0.061)	0.197*** (0.064)
STFT	0.337*** (0.077)	0.196*** (0.062)
Grim	0.062 (0.043)	0.193*** (0.066)
Grim2	0.033 (0.031)	0.061 (0.047)
Grim3		0.077 (0.049)
T2		0.043 (0.033)
T3	0.030 (0.029)	
β	0.892 (0.018)	0.882 (0.017)
LL		-1077.367
BIC		2227.427

Table 3 Estimated shares of strategies selected by the Bayesian Information Criterion (BIC). Shares with an estimated value of zero are not displayed to improve readability. Standard errors are bootstrapped from 10,000 resamples and reported in parentheses. The statistical significance of each share is assessed based on the estimated distribution of the parameter. β is implied by the estimated tremble and represents the probability that subjects follow what is prescribed by the strategies. The log-likelihood (LL) and the Bayesian Information Criterion (BIC) of the joint model fit are displayed in the last two rows. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

The table shows a clear composition shift under the *Tournament* treatment. Relative to *Control*, the competitive environment is associated with a pronounced reallocation of mass away from strategies that start by defecting toward trigger-type punishment, as Grim is relevant only in the *Tournament* condition and equal to 0.19. The parameter β , which is implied by the estimated tremble, is similar across treatments ($\beta_C = 0.892$, $\beta_T = 0.882$), indicating that the treatment effect operates through strategy composition rather than decision noise.

Taken together, these estimates indicate that intergroup competition does not simply raise cooperation, but it changes how cooperation is sustained. In the *Control* condition, a large share of subjects rely on AD and STFT, both of which start by defecting, with STFT allowing for some recovery after a deviation. Under *Tournament*, their combined weight falls, while the share of Grim, and to a lesser extent, TFT, increases. Grim is less forgiving than STFT, as any defection permanently triggers punishment. Thus, the rise in cooperation documented above is not driven by greater leniency after deviations, but by a shift toward strategies that begin with cooperation and impose stricter sanctions when

cooperation is broken, consistent with the idea that the tournament strengthens a social norm to cooperate. In this sense, intergroup competition both enhances initial cooperation and reallocates probability mass toward trigger strategies that make cooperative play more sustainable, in line with the mechanism formalized in Section 4, where the tournament shifts equilibrium selection toward cooperative trigger strategies.

4 Theoretical Framework

The results presented in the previous section indicate that intergroup competition increases cooperation and alters the distribution of strategies: fewer participants choose Always Defect, while a substantial share adopts the trigger strategy Grim. In this section, we provide a theoretical analysis of the game used in the experiment. Specifically, we investigate how competition affects the equilibria of the game by introducing a tournament between pairs of players who play an infinitely repeated Prisoner’s Dilemma. The pair that achieves the highest cumulative sum of aggregate payoffs (points) wins the tournament. Because cooperation yields more points, cooperating increases the odds of winning. No additional monetary payoffs are awarded to the winners.

Intuitively, the tournament introduces an additional incentive to cooperate. This extra incentive modifies the strategic environment and lowers the thresholds under which cooperative trigger strategies are viable. In particular, we derive conditions under which the tournament reduces both the subgame-perfect and the risk-dominance thresholds. The model therefore predicts a reallocation of play away from Always Defect toward the least risky cooperative strategy, Grim, a pattern that closely matches what we observe in the data.

To obtain these results, we model the incentives introduced by the tournament as hedonic utility from winning. This choice is guided by two considerations. First, it is parsimonious, as we simply posit that players enjoy winning the tournament and that cooperation increases the likelihood of doing so. Second, the specification keeps agents selfish with respect to material payoffs and does not require the presence or activation of other-regarding preferences.

It is important to note that similar comparative statics to those in Propositions 1 and 2 could be obtained in alternative frameworks.¹¹ For instance, group-identity social preferences (Y. Chen & Li, 2009) or social framing (Tversky & Kahneman, 1981) could also increase the relative value of mutual cooperation and thereby lower the critical discount factors for cooperative equilibria. Our aim in this section is therefore not to test the

¹¹For example, B. Chen et al. (2025) incorporates social preferences and group identity into an infinitely repeated Prisoner’s Dilemma and shows that these factors reduce both the subgame-perfect discount-factor threshold and the risk-dominance cutoff for Grim relative to Always Defect, yielding comparative statics that are closely related to ours.

specific hypothesis that players derive hedonic utility from winning. Rather, we use this simple tournament-based model to characterize how an intergroup competition without material prizes can change equilibrium selection in the repeated Prisoner’s Dilemma and to provide a theoretical benchmark that is consistent with the observed shift from Always Defect to Grim.

4.1 The Model

In what follows, we analyze the game described in the previous section, an indefinitely repeated PD. To ease the exposition, we perform a normalization of the payoff matrix as shown in Table 4.¹² To ensure that mutual cooperation generates a higher combined outcome, it is required that $2 > 1 + g - \ell > 0$. Otherwise, alternating between $(Cooperate, Defect)$ and $(Defect, Cooperate)$ would generate higher payoffs for both players. This condition, which is satisfied in our experimental design, ensures that the benefits of cooperation outweigh the potential gains from alternating between cooperation and defection.

	<i>Cooperate</i>	<i>Defect</i>
<i>Cooperate</i>	1	$-\ell$
<i>Defect</i>	$1 + g$	0

Table 4 Row Player’s Payoffs of the stage game.

The first step is to compute, from the game parameters, the thresholds δ^{SPE} and δ^{RD} . To sustain cooperation in a subgame perfect equilibrium (SPE), players must be sufficiently forward-looking, meaning that the discount factor δ^{SPE} must be large enough. Specifically, cooperation can be sustained if the discount factor satisfies the following condition:

$$\delta \geq \delta^{SPE} = \frac{g}{1 + g}.$$

We can also determine the minimum value of δ required for cooperation to be part of a risk-dominant equilibrium.¹³ This condition is met when:

$$\delta \geq \delta^{RD} = \frac{g + \ell}{1 + g + \ell}.$$

The Tournament. In the tournament setting, two pairs of players, referred to as teams, engage in an infinitely repeated PD game. Each player is aware of the presence

¹²We performed the same normalization as in Dal Bó and Fréchette (2018). For our game, the game’s parameters are set to: $g = \frac{18}{7}$ and $\ell = \frac{13}{7}$, while the continuation probability is set to $\delta = 0, 75$.

¹³Harsanyi and Selten (1988) define risk dominance for 2×2 games. It is possible to extend the concept of risk dominance to repeated games using auxiliary 2×2 games that implement specific equilibrium strategies. For more reference, see Blonski and Spagnolo (2015).

of the opposing team. The objective of the tournament is for a team to achieve the highest number of *points* (aggregate sum of both players' individual payoffs). Winning the tournament does not provide any additional material payoff. Furthermore, the actions taken by one team do not directly affect the payoffs of the other team, and vice versa.

Assume that each player experiences a non-negative hedonic utility, denoted as $W \geq 0$, when their team wins the tournament. This utility is in addition to the monetary payoffs normally obtained from the game.

In our setting, the probability that a team wins the tournament, denoted $\mathbb{P}(\text{win} | s)$, depends on the total number of *points* the team expects to accumulate under the strategy profile s . Since the tournament rule is that the team with the highest score wins, this probability increases with the expected *points* scored by the team following the prescribed strategy. As a result, when choosing a strategy, players must weigh not only the immediate payoffs from their actions, but also how these actions affect their team's chances of winning, and thus the additional utility W they derive from victory. In the following, we focus on three benchmark strategies: Always Cooperate (AC), Grim (G), and Always Defect (AD). Given the structure of the stage-game payoffs, with $2 > 1 + g - \ell > 0$, it follows that AC yields the highest expected score, followed by G and then AD. Consequently $\mathbb{P}(\text{win} | \text{AC}) \geq \mathbb{P}(\text{win} | \text{G}) \geq \mathbb{P}(\text{win} | \text{AD})$.

This ranking suffices to establish our main results; no additional assumptions on the functional form of $\mathbb{P}(\text{win} | s)$ are required. Although the probability of victory depends on the strategies chosen by both teams, a given team cannot control its opponent's behavior. Accordingly, we treat $\mathbb{P}(\text{win} | s)$ as parametrically fixed from each team's viewpoint and analyze how a team's own strategy affects its expected outcome. In this sense, the results are derived in partial equilibrium. Solving for a general equilibrium that simultaneously determines both teams' strategies would demand stronger assumptions about belief formation and strategic interdependence; in our opinion such an extension is unnecessary here, because the ordinal ranking of strategies, together with the monotonic relationship between *points* and winning, is sufficient to show how the tournament incentives lower the threshold for cooperation.

To prove that the tournament lowers the threshold δ^{SPE} necessary for cooperation, we follow the steps of Nash reversion and incorporate into the payoff of each strategy the value of winning the tournament, W , weighted by the probability of winning given the strategy played. This leads us to the first result:

Proposition 1. *Let W be the utility given by winning the tournament, then the minimum discount factor necessary to have cooperation as part of an SPE in the presence of a tournament, δ^{SPE*} , is lower than δ^{SPE} in the absence of the tournament. Moreover δ^{SPE*}*

is equal to:

$$\delta^{SPE^*} = \frac{g - W(\mathbb{P}(\text{win} | AC) - \mathbb{P}(\text{win} | G))}{1 + g - W(\mathbb{P}(\text{win} | AC) - \mathbb{P}(\text{win} | G))} \leq \frac{g}{1 + g} = \delta^{SPE}.$$

In particular, $\mathbb{P}(\text{win} | AC)$ denotes the probability of winning the tournament when both players continue to cooperate from period t^* onward, whereas $\mathbb{P}(\text{win} | G)$ denotes the probability of winning along the grim punishment path that follows a one-shot deviation at t^* (see Appendix A). Under this interpretation, $\mathbb{P}(\text{win} | AC) - \mathbb{P}(\text{win} | G) > 0$, which yields the expression for δ^{SPE^*} above.

This first result demonstrates that the tournament reduces the threshold for cooperation to be sustained in an SPE. This implies that competition can enhance cooperation, even without altering the stage game's payoffs.

To prove that the tournament lowers the threshold for a risk-dominant equilibrium δ^{RD} , we follow Blonski and Spagnolo (2015). To determine when cooperation is risk-dominant, we focus exclusively on two equilibria in pure actions: the Grim strategy , which is the least risky among cooperative equilibria, and always defect (AD).¹⁴ By following the steps outlined in Blonski and Spagnolo (2015), we derive the following result:

Proposition 2. *Let W be the utility given by winning the tournament, then the minimum discount factor necessary to have cooperation as part of a risk-dominant strategy in the presence of a tournament, δ^{RD^*} , is lower than δ^{RD} in the absence of the tournament. Moreover, δ^{RD^*} is equal to:*

$$\delta^{RD^*} = \frac{g + \ell - W(\mathbb{P}(\text{win} | AC) - \mathbb{P}(\text{win} | AD))}{1 + g + \ell - W(\mathbb{P}(\text{win} | AC) - \mathbb{P}(\text{win} | AD))} \leq \frac{g + \ell}{1 + g + \ell} = \delta^{RD}.$$

This second result is particularly relevant in our setting because, in our game, the parameters are such that cooperation can be sustained in an SPE, but it is not risk dominant. The game parameters used in our experiment are such that cooperation can be sustained in equilibrium, as the derived δ^{SPE} is 0.72, which is lower than the continuation probability of $\delta = 0.75$. However, cooperation is not risk-dominant, as indicated by the fact that $\delta^{RD} \approx 0.82$ is higher than the continuation probability. As Proposition 2 suggests, players who care enough about winning might switch from Always Defect to Grim in the *Tournament* treatment. This is actually the result we empirically observe, which we report in Section 3.2, where we estimate the strategies played by participants. The detailed proofs for both propositions can be found in appendix A.

¹⁴Proof in Blonski and Spagnolo (2015).

5 Conclusions

In this study, we provide evidence that competition between groups, implemented with a tournament, can substantially increase cooperation in strategic environments, even when it does not confer material rewards to the winners. The experimental results show higher cooperation rates in the tournament than in the control condition, with the effect strengthening as participants gain experience. We also uncover how competition shapes strategic behavior, promoting a shift away from strategies that prescribe starting with defection toward the least risky cooperative strategy, Grim. This insight relies on the features of the Prisoner’s Dilemma game, which represents a key methodological innovation compared to the existing literature.

To interpret these patterns, we develop a simple tournament-based model of the infinitely repeated Prisoner’s Dilemma. The intergroup contest adds a non-monetary incentive to win, which lowers both the subgame-perfect and risk-dominance discount-factor thresholds for Grim relative to Always Defect. The model thus predicts the observed empirical reallocation from Always Defect to Grim and provides a tractable benchmark for how non-monetary competition can change equilibrium selection in repeated interactions.

These findings highlight the effectiveness of non-monetary competition in fostering cooperation and underline the added value of our methodological approach. The use of the indefinitely repeated Prisoner’s Dilemma enabled us to estimate the strategies employed by participants and to connect the experimental evidence to a tractable theoretical framework. Taken together, the empirical and theoretical analyses provide a coherent account of how intergroup competition can discipline and coordinate behavior, shifting play toward cooperative trigger strategies even in the absence of material prizes.

References

- Abbink, K., Brandts, J., Herrmann, B., & Orzen, H. (2010). Intergroup conflict and intra-group punishment in an experimental contest game. *The American Economic Review*, 100(1), 420–447.
- Alchian, A. A., & Demsetz, H. (1972). Production, information costs, and economic organization. *The American Economic Review*, 62(5), 777–795.
- Aoyagi, M., Bhaskar, V., & Fréchette, G. R. (2019, February). The impact of monitoring in infinitely repeated games: Perfect, public, and private. *American Economic Journal: Microeconomics*, 11(1), 1–43.
- Aoyagi, M., Fréchette, G. R., & Yuksel, S. (2024, December). Beliefs in repeated games: An experiment. *American Economic Review*, 114(12), 3944–75.
- Arechar, A. A., Dreber, A., Fudenberg, D., & Rand, D. G. (2017). “i’m just a soul whose intentions are good”: The role of communication in noisy repeated games. *Games*

- and Economic Behavior*, 104, 726–743.
- Augenblick, N., & Cunha, J. M. (2015). Competition and cooperation in a public goods game: A field experiment. *Economic inquiry*, 53(1), 574–588.
- Bacharach, M. (1999). Interactive team reasoning: A contribution to the theory of co-operation. *Research in economics*, 53(2), 117–147.
- Bigoni, M., Casari, M., Skrzypacz, A., & Spagnolo, G. (2015). Time horizon and cooperation in continuous time. *Econometrica*, 83(2), 587–616.
- Blonski, M., & Spagnolo, G. (2015). Prisoners' other dilemma. *International Journal of Game Theory*, 44, 61–81.
- Breitmoser, Y. (2015). Cooperation, but no reciprocity: Individual strategies in the repeated prisoner's dilemma. *American Economic Review*, 105(9), 2882–2910.
- Burton-Chellew, M. N., Ross-Gillespie, A., & West, S. A. (2010). Cooperation in humans: competition between groups and proximate emotions. *Evolution and Human behavior*, 31(2), 104–108.
- Camera, G., Casari, M., & Bigoni, M. (2012). Cooperative strategies in anonymous economies: An experiment. *Games and Economic Behavior*, 75(2), 570–586.
- Cárdenas, J. C., & Mantilla, C. (2015). Between-group competition, intra-group cooperation and relative performance. *Frontiers in behavioral neuroscience*, 9, 33.
- Cason, T. N., Lau, S.-H. P., & Mui, V.-L. (2019). Prior interaction, identity, and cooperation in the inter-group prisoner's dilemma. *Journal of Economic Behavior & Organization*, 166, 613–629.
- Cheikbossian, G. (2012). The collective action problem: Within-group cooperation and between-group competition in a repeated rent-seeking game. *Games and Economic Behavior*, 74(1), 68–82.
- Chen, B., Li, S. X., & Liu, T. X. (2025). *Theory appendix for “group identity and cooperation in infinitely repeated games”*. Retrieved from https://people.smu.edu/bchen/files/2025/05/Theory_Appendix.pdf (Working paper)
- Chen, Y., & Li, S. X. (2009, March). Group identity and social preferences. *American Economic Review*, 99(1), 431–57.
- Chen, Y.-Y. (2020). Intergroup competition with an endogenously determined prize level. *Journal of Economic Behavior & Organization*, 178, 759–776.
- Dal Bó, P., & Fréchette, G. R. (2011). The evolution of cooperation in infinitely repeated games: Experimental evidence. *American Economic Review*, 101(1), 411–429.
- Dal Bó, P., & Fréchette, G. R. (2018). On the determinants of cooperation in infinitely repeated games: A survey. *Journal of Economic Literature*, 56(1), 60–114.
- Dvorak, F. (2023). stratest: a software package for strategy frequency estimation. *Journal of the Economic Science Association*, 9(2), 337–349.
- Ellingsen, T., Johannesson, M., Mollerstrom, J., & Munkhammar, S. (2012). Social framing effects: Preferences or beliefs? *Games and Economic Behavior*, 76(1),

117–130.

- Fehr, E., & Gächter, S. (2000). Cooperation and punishment in public goods experiments. *American Economic Review*, 90(4), 980–994.
- Fischbacher, U. (2007). z-tree: Zurich toolbox for ready-made economic experiments. *Experimental Economics*, 10, 171–178.
- Fréchette, G. R., & Yuksel, S. (2017). Infinitely repeated games in the laboratory: Four perspectives on discounting and random termination. *Experimental Economics*, 20, 279–308.
- Fudenberg, D., & Maskin, E. (1986). The folk theorem in repeated games with discounting or with incomplete information. *Econometrica*, 54(3), 533–554.
- Fudenberg, D., Rand, D. G., & Dreber, A. (2012, April). Slow to anger and fast to forgive: Cooperation in an uncertain world. *American Economic Review*, 102(2), 720–49.
- Greiner, B. (2015). Subject pool recruitment procedures: organizing experiments with orsee. *Journal of the Economic Science Association*, 1(1), 114–125.
- Hamilton, B. H., Nickerson, J. A., & Owan, H. (2003). Team incentives and worker heterogeneity: An empirical analysis of the impact of teams on productivity and participation. *Journal of Political Economy*, 111(3), 465–497.
- Harsanyi, J. C., & Selten, R. (1988). A general theory of equilibrium selection in games. *MIT Press Books*, 1.
- Holmstrom, B. (1982). Moral hazard in teams. *The Bell Journal of Economics*, 13(2), 324–340.
- Jones, M. T. (2014). Strategic complexity and cooperation: An experimental study. *Journal of Economic Behavior & Organization*, 106, 352–366.
- Markussen, T., Reuben, E., & Tyran, J.-R. (2014). Competition, cooperation and collective choice. *The Economic Journal*, 124(574), F163–F195.
- Ostrom, E. (1990). *Governing the commons: The evolution of institutions for collective action*. Cambridge University Press.
- Puurtinen, M., & Mappes, T. (2009). Between-group competition and human cooperation. *Proceedings of the Royal Society B: Biological Sciences*, 276(1655), 355–360.
- Reuben, E., & Tyran, J.-R. (2010). Everyone is a winner: Promoting cooperation through all-can-win intergroup competition. *European Journal of Political Economy*, 26(1), 25–35.
- Romero, J., & Rosokha, Y. (2023). Mixed strategies in the indefinitely repeated prisoner's dilemma. *Econometrica*, 91(6), 2295–2331.
- Schnieder, C. (2022). How relative performance information affects employee behavior: a systematic review of empirical research. *Journal of Accounting Literature*, 44(1), 72–107.
- Tan, J. H., & Bolle, F. (2007). Team competition and the public goods game. *Economics Letters*, 96(1), 133–139.

- Tversky, A., & Kahneman, D. (1981). The framing of decisions and the psychology of choice. *Science*, 211(4481), 453–458.
- Vespa, E. (2020). An experimental investigation of cooperation in the dynamic common pool game. *International Economic Review*, 61(1), 417–440.

Appendix for “The Enemy of My Enemy: How Competition Mitigates Social Dilemmas”

Alessandro Strighi, Sara Gil-Gallen, Andrea Albertazzi

A Proofs

Proof of Proposition 1 In order to prove Proposition 1, we follow the proof of Nash reversion, and we add to each strategy the value of winning the tournament W weighted by the probability of winning given the strategy played. We compare the continuation values from period t^* onward between (i) continuing with AC and (ii) deviating once at t^* and then following the grim punishment path G . Therefore, the equation becomes the following:

$$\sum_{t=t^*}^{\infty} \delta^t \cdot 1 + W\mathbb{P}(\text{win} \mid AC) \geq 1 + g + \sum_{t=t^*+1}^{\infty} \delta^t \cdot 0 + W\mathbb{P}(\text{win} \mid G)$$

where AC is the continuation strategy in which both players keep cooperating, while G is the Grim strategy in which the player “pulls the trigger” at time t^* , and after that, both players play *Defect*. Since $2 > 1 + g - \ell$, the strategy AC gives more points than the strategy G , thus $\mathbb{P}(\text{win} \mid AC) - \mathbb{P}(\text{win} \mid G) \geq 0$.

Rearranging the formula, we obtain:

$$\delta^{SPE^*} = \frac{g - W(\mathbb{P}(\text{win} \mid AC) - \mathbb{P}(\text{win} \mid G))}{1 + g - W(\mathbb{P}(\text{win} \mid AC) - \mathbb{P}(\text{win} \mid G))} \leq \frac{g}{1 + g}, \quad \forall W > 0.$$

□

Proof of Proposition 2 In order to prove proposition 2, we follow Blonski and Spagnolo (2015). To assess when coordination is risk-dominant, we focus only on two equilibria in pure actions: the grim trigger strategy (G), which is the least risky among cooperative equilibria (proof in Blonski and Spagnolo (2015)), and always defect (AD). We build an accessory 2×2 game using only these two equilibrium points. According to Harsanyi and Selten (1988), risk dominance in 2×2 games can be determined by comparing the Nash-products of the two equilibria, namely the product of both players’ disincentives not to behave according to the equilibrium under consideration. We evaluate all continuation payoffs from period t^* onward. Factoring out the common discount term δ^{t^*} , we can write the disincentives u_i for G and v_i for AD as:

$$u_i = \sum_{t=t^*}^{\infty} \delta^t \cdot 1 - \left(1 + g + \sum_{t=t^*+1}^{\infty} \delta^t \cdot 0 \right) \geq 0$$

$$v_i = \sum_{t=t^*}^{\infty} \delta^t \cdot 0 - \left(-l + \sum_{t=t^*+1}^{\infty} \delta^t \cdot 0 \right) \geq 0.$$

The grim trigger strategy G is risk dominated by AD if $v_1 v_2 \geq u_1 u_2$:

$$\ell^2 - \left(\frac{1}{1-\delta} - (1+g) \right)^2 \geq 0.$$

From these relations, we find that the threshold for δ below which G is risk-dominated is the following:

$$\delta^{RD} = \frac{g+l}{1+g+l}.$$

Similarly to proposition 1, we add the weighted value of winning the tournament. Therefore, the relations become:

$$u_i = \sum_{t=t^*}^{\infty} \delta^t \cdot 1 + W\mathbb{P}(\text{win} \mid \text{AC}) - \left(1 + g + \sum_{t=t^*+1}^{\infty} \delta^t \cdot 0 + W\mathbb{P}(\text{win} \mid \text{G}) \right) \geq 0$$

$$v_i = \sum_{t=t^*}^{\infty} \delta^t \cdot 0 + W\mathbb{P}(\text{win} \mid \text{AD}) - \left(-l + \sum_{t=t^*+1}^{\infty} \delta^t \cdot 0 + W\mathbb{P}(\text{win} \mid \text{G}) \right) \geq 0.$$

Using the same procedures as before, we obtain,

$$\left(l + W \left(\mathbb{P}(\text{win} \mid \text{AD}) - \mathbb{P}(\text{win} \mid \text{G}) \right) \right)^2 - \left(\frac{1}{1-\delta} - (1+g) + W \left(\mathbb{P}(\text{win} \mid \text{AC}) - \mathbb{P}(\text{win} \mid \text{G}) \right) \right)^2 \geq 0$$

and by rearranging the formula, we obtain:

$$\delta^{RD*} = \frac{g+l - W \left(\mathbb{P}(\text{win} \mid \text{AC}) - \mathbb{P}(\text{win} \mid \text{AD}) \right)}{1+g+l - W \left(\mathbb{P}(\text{win} \mid \text{AC}) - \mathbb{P}(\text{win} \mid \text{AD}) \right)} \leq \frac{g+\ell}{1+g+\ell}, \quad \forall W > 0.$$

Where the strategies AC, AD, and G are, respectively, Always Cooperate, Always Defect, and the Grim strategy. \square

B Instructions

B.1 Control Treatment

Welcome

You are about to participate in a session on decision-making, and you will be paid for your participation with cash vouchers, privately, at the end of the session. What you earn depends partly on your decisions, partly on the decisions of others, and partly on chance.

General Instructions

1. In this experiment, you will be asked to make decisions in several rounds. You will be randomly paired with another person for a sequence of rounds. Each sequence of rounds is referred to as a match.
2. The length of a match is randomly determined. After each round, there is a 75% probability that the match will continue for at least another round. This probability is always the same regardless of the round. So, for instance, if you are in round 2, the probability there will be a third round is 75%, and if you are in round 9, the probability there will be another round is also 75%.
3. At the beginning of a new match, you will be randomly paired with another person for a new match.
4. The choices and the payoffs (expressed in points) in each round are as follows:

		The other's choice	
		1	2
Your choice	1	(32 , 32)	(12 , 50)
1			
2		(50 , 12)	(25 , 25)

The first entry in each cell represents your payoff, while the second entry represents the payoff of the person you are matched with.

For example, if:

- You select **1** and the other selects **1**, you each make 32.
- You select **1** and the other selects **2**, you make 12 while the other makes 50.
- You select **1** and the other selects **2**, you make 50 while the other makes 12.
- You select **2** and the other selects **2**, you each make 25.

5. At the end of the 50 min, you will be paid 0.005€ (half of a euro cent) for every point you scored individually in every round played during the whole experiment.
6. Are there any questions?

B.2 Tournament treatment

All the framing introduced in the instructions for the treatment that does not appear in control is indicated in italics.

Welcome

You are about to participate in a session on a tournament, and you will be paid for your participation with cash vouchers, privately at the end of the session. What you earn depends partly on your decisions, partly on the decisions of others, and partly on chance.

General Instructions

1. In this experiment, you will be asked to make decisions in several rounds. You will be randomly paired with *a teammate* for a sequence of rounds. Each sequence of rounds is referred to as a match.
2. *During each match, your team will compete against one adversary team randomly chosen between the other teams in this experiment. The team that earns more points at the end of the match will be declared the winner.*
3. The length of a match is randomly determined. After each round, there is a 75% probability that the match will continue for at least another round. This probability is always the same regardless of the round. So, for instance, if you are in round 2, the probability there will be a third round is 75%, and if you are in round 9, the probability there will be another round is also 75%. *The match will end for both teams at the same time.*
4. At the beginning of a new match, you will be randomly paired with another *teammate*, *and you will play against a new adversary team.*
5. The choices and the payoffs (expressed in points) in each round are as follows:

Teammate's choice		
Your choice	1	2
1	(32 , 32)	(12 , 50)
2	(50 , 12)	(25 , 25)

The first entry in each cell represents your payoff, while the second entry represents the payoff of your *teammate*. *The sum of your payoff and your teammate's payoff in each round during the whole match will determine your total team's points in the match.*

For example, if:

- You select **1** and the *teammate* selects **1**, you each make 32. *The team's points in the round will be equal to 64.*
- You select **1** and the *teammate* selects **2**, you make 12 while the *teammate* makes 50. *The team's points in the round will be equal to 62.*
- You select **2** and the *teammate* selects **1**, you make 50 while the *teammate* makes 12. *The team's points in the round will be equal to 62.*
- You select **2** and the *teammate* selects **2**, you each make 25. *The team's points in the round will be equal to 50.*

If the total points of your team are higher than the total points of the adversary team, your team wins the match, otherwise, your team loses.

6. At the end of the 50 min, you will be paid 0.005€ (half of a euro cent) for every point you scored individually in every round played during the whole experiment.
Note that you will not earn any additional money for winning a match.
7. Are there any questions?

B.3 Questionnaire

Socio-Demographics

- How old are you?
- What is your gender? Male Female
- What is your occupation?
 - Student
 - Employee
 - Unemployed
 - Retired
 - Other
- What is your field of study?
 - Economics and management
 - Social Sciences
 - Arts and Humanities
 - Engineering Sciences
 - Medical studies
 - Other
- How much experience have you had with LEEN before?

Psychological questions

- From 0 to 10, how much do you trust people in general, where 0 indicates “better not trust none” and 10 means “better completely trust”?

0 1 2 3 4 5 6 7 8 9 10

- For a scale from 0 to 10, how do you evaluate your behavior in front of risk: you are a person who avoids risk (1), or do you love risk (10)?

0 1 2 3 4 5 6 7 8 9 10

C Sample

Table C.1 reports the results of OLS regressions of *Tournament* on the relevant variable elicited in the questionnaire. The estimates show the treatment assignment was balanced with respect to all these variables.

			Economic	Lab			
	Age	Female	Student	background	experience	Risk	Trust
<i>Tournament</i>	-1.042 (1.05)	-0.111 (0.10)	0.049 (0.07)	0.155 (0.10)	0.063 (0.57)	0.011 (0.41)	-0.096 (0.38)
Constant	24.5*** (0.75)	0.674*** (0.07)	0.826*** (0.05)	0.283*** (0.07)	2.978*** (0.41)	5.739*** (0.29)	5.804*** (0.27)
Participants	94	94	94	94	94	94	94
R-squared	0.011	0.013	0.005	0.026	0.000	0.000	0.001

Table C.1 Balancing test. Standard errors in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

D Strategies Estimation

The following are the 17 strategies we consider in our initial set of strategies. Always Defect (AD), Always Cooperate (AC), Grim, Tit for Tat (TFT), 2TFT defects if opponent defected in either of the last 2 rounds, Win Stay Lose Shift (WSLS), T2-T8 are threshold strategies that cooperate until round 2-8 and then defects forever, suspicious Tit for Tat (STFT) is equal to TFT with the only difference being it starts by defecting, Grim2-3 are lenient version of Grim that defect forever following a defection of the opponent in the last 2-3 rounds, and Tit for Two Tat (TF2T), a lenient version of TFT that retaliates only after an opponent has defected twice in a row.

D.1 Decisions from supergames 20-29

	<i>Control</i>	<i>Tournament</i>
AD	0.378*** (0.078)	0.233*** (0.065)
AC		
TFT	0.130** (0.056)	0.197*** (0.064)
STFT	0.323*** (0.076)	0.196*** (0.062)
Grim	0.052 (0.038)	0.191*** (0.065)
Grim2		0.053 (0.041)
Grim3		0.052 (0.044)
TF2T	0.025 (0.026)	0.035 (0.042)
2TFT	0.031 (0.032)	
WSLS		
T2		0.043 (0.032)
T3	0.022 (0.021)	
T4	0.022 (0.022)	
T5		
T6	0.017 (0.021)	
T7		
T8		
β	0.896 (0.017)	0.883 (0.017)
LL		-1072.591
BIC		2299.653

Table D.1 Estimated shares of strategies from the full initial set. Shares with an estimated value of zero are not displayed to improve readability. Standard errors are bootstrapped from 10,000 resamples and reported in parentheses. The statistical significance of each share is assessed based on the estimated distribution of the parameter. β is implied by the estimated tremble and represents the probability that subjects follow what is prescribed by the strategies. The log-likelihood (LL) and the Bayesian Information Criterion (BIC) of the joint model fit are displayed in the last two rows. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

	Shares		β	
	<i>Control</i>	<i>Tournament</i>	<i>Control</i>	<i>Tournament</i>
AD	0.305*** (0.069)	0.194** (0.091)	0.988 (0.006)	0.909 (0.026)
TFT	0.103** (0.050)	0.155*** (0.079)	0.936 (0.054)	0.955 (0.056)
STFT	0.432*** (0.078)	0.195*** (0.101)	0.812 (0.024)	0.879 (0.055)
Grim	0.059 (0.039)	0.209** (0.087)	0.98 (0.054)	0.91 (0.030)
Grim2	0.022 (0.022)	0.041 (0.029)	0.000 (0.147)	0.000 (0.065)
Grim3		0.089 (0.064)		0.836 (0.062)
T2	0.031 (0.030)	0.021 (0.029)	0.398 (0.078)	0.000 (0.229)
T3	0.047 (0.034)	0.095 (0.058)	0.727 (0.119)	0.556 (0.190)
LL		-990.3886		
BIC		2117.076		

Table D.2 Estimated shares of strategies selected by the Bayesian Information Criterion (BIC) with different trembles per strategy. Shares with an estimated value of zero are not displayed to improve readability. Standard errors are bootstrapped from 10,000 resamples and reported in parentheses. The statistical significance of each share is assessed based on the estimated distribution of the parameter. β is implied by the estimated tremble and represents the probability that subjects follow what is prescribed by the strategies. The log-likelihood (LL) and the Bayesian Information Criterion (BIC) of the joint model fit are displayed in the last two rows. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

D.2 Decisions from all supergames

	<i>Control</i>	<i>Tournament</i>
AD	0.382*** (0.075)	0.228*** (0.063)
TFT	0.240*** (0.064)	0.197*** (0.061)
STFT	0.241*** (0.067)	0.184*** (0.059)
Grim	0.111** (0.051)	0.300*** (0.068)
Grim2	0.026 (0.026)	0.021 (0.021)
2TFT		0.049 (0.034)
T8		0.021 (0.021)
β	0.844 (0.016)	0.838 (0.014)
LL		-4623.455
BIC		9310.516

Table D.3 Estimated shares of strategies selected by the Bayesian Information Criterion (BIC) using all available decisions. Shares with an estimated value of zero are not displayed to improve readability. Standard errors are bootstrapped from 10,000 resamples and reported in parentheses. The statistical significance of each share is assessed based on the estimated distribution of the parameter. β is implied by the estimated tremble and represents the probability that subjects follow what is prescribed by the strategies. The log-likelihood (LL) and the Bayesian Information Criterion (BIC) of the joint model fit are displayed in the last two rows. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

	<i>Control</i>	<i>Tournament</i>
AD	0.382*** (0.075)	0.229*** (0.064)
AC		
TFT	0.240*** (0.065)	0.172*** (0.058)
STFT	0.242*** (0.066)	0.185*** (0.060)
Grim	0.095** (0.047)	0.273*** (0.068)
Grim2	0.023 (0.022)	0.021 (0.021)
Grim3		
TF2T		0.023 (0.024)
twoTFT		0.05 (0.035)
WSLS		
T2		0.027 (0.027)
T3	0.019 (0.021)	
T4		
T5		
T6		
T7		
T8		0.021 (0.021)
β	0.844 (0.016)	0.84 (0.014)
LL		-4617.302
BIC		9389.077

Table D.4 Estimated shares from the full set of strategies using all available decisions. Shares with an estimated value of zero are not displayed to improve readability. Standard errors are bootstrapped from 10,000 resamples and reported in parentheses. The statistical significance of each share is assessed based on the estimated distribution of the parameter. β is implied by the estimated tremble and represents the probability that subjects follow what is prescribed by the strategies. The log-likelihood (LL) and the Bayesian Information Criterion (BIC) of the joint model fit are displayed in the last two rows. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$